


1 Introduction


1.1 What is a Linked List

A Linked List is a simple data structure consisting of what is usually referred to as *Nodes*. A Node contains some piece of data as well as a pointer to the next node (it can also contain a pointer to the previous node, to create a Doubly Linked List, which is actually what is used in the implementation of SqrtList). Thus, the whole list can be accessed by starting at the first node (usually referred to as the *head*) and following the pointers until the end of the list (*tail*) is reached. A Linked List does not have random access, and thus although given a specific node all operations (inserting a new node next to it, removing the node, getting its value or setting its value) are done in $O(1)$ time, doing the same operations given an *index* instead of a node are all done in $O(n)$ time (where n is the length of the list), as, in the worst case, the entire list needs to be traversed to access the node at that index and perform the operation.

A Singly Linked List




A Doubly Linked List



1.2 What is the Square Root List

The *Square Root List* is an idea I had after examining Linked Lists in my Data Structures class and it is an attempt to improve on the runtime of the aforementioned operations by storing some extra pointers along with the real list. A second list is maintained, which stores pointers for about every \sqrt{n} elements of the main list, where n is the length of the main list. When an element of the main list needs to be accessed, instead of traversing the main list directly, the appropriate element is selected from the secondary list as the traversal starting point, thus "skipping" over many elements; then the traversal of the main list starts and the element is reached. We will show later that this traversal is done in $O(\sqrt{n})$ time.

A Square Root List



2 Operation Runtime for the Square Root List

2.1 Notes

Before showing that the runtime is indeed what was claimed before, a point needs to be clarified. I said that the secondary list is about \sqrt{n} elements long. That "about" means the following: because add and remove operations are constantly changing the length n of the list, instead of constantly keeping up with those changes, the list uses another number m as a stand-in for the length of the list. m is updated in the `buildMeta()` method of the class, where it is set to be equal to the real size of the list. `buildMeta` is only called when the size of the list has diverged too much from the value of m . When an insertion happens but the threshold for calling `buildMeta` is not surpassed, then `retreat` is called. `retreat` traverses the secondary list from the site of the insertion until the end and updates each stored node to be equal to its next node. This ensures that the chunks between two stored nodes remain equally sized and that the list works correctly. Similarly, when a removal happens, `advance` is called which works identically to `retreat`, but updates nodes to be their previous node instead of their next.

2.2 Proofs

Claim: \sqrt{m} is bounded by $O(\sqrt{n})$.

Proof.

We can see from the code that $m - \sqrt{m} \leq n \leq m + \sqrt{m}$, since when either of those bounds is exceeded, `buildMeta` is called and then $m = n$.

$m - \sqrt{m} \leq n \iff m - \sqrt{m} - n \leq 0$ (inequality (1))

By setting $x = \sqrt{m}$ inequality (1) becomes $x^2 - x - n \leq 0$ (inequality (2))

$x^2 - x - n = 0 \iff x = \frac{1+\sqrt{1+4n}}{2}$ or $x = \frac{1-\sqrt{1+4n}}{2}$

Therefore inequality (2) is satisfied by $\frac{1-\sqrt{1+4n}}{2} \leq x \leq \frac{1+\sqrt{1+4n}}{2}$.

This gives us that $0 \leq \sqrt{m} \leq \frac{1+\sqrt{1+4n}}{2}$.

$\sqrt{m} \leq \frac{1}{2} + \frac{\sqrt{1+4n}}{2} \leq \frac{\sqrt{n}}{2} + \frac{\sqrt{5n}}{2}$

Thus $\sqrt{m} \leq \frac{\sqrt{5}+1}{2}\sqrt{n}$ and therefore \sqrt{m} is bounded by $O(\sqrt{n})$. □

Claim: The `getNode(idx)` method works in $O(\sqrt{n})$ time.

Proof.

The appropriate starting node from the secondary list is selected in constant time. Then, a traversal is needed, from the node stored in the secondary list up to the desired index. That traversal has at most \sqrt{m} nodes, since, the secondary list stores one node every \sqrt{m} nodes. Thus, the entire traversal is at most $\sqrt{m} - 1$ nodes long, which means that it is bounded by $O(\sqrt{m})$ and thus $O(\sqrt{n})$. □

Claim: The `get(idx)` method works in $O(\sqrt{n})$ time.

Proof.

The method consists of a single call to `getNode(idx)` a field access. Since `getNode(idx)` works in $O(\sqrt{n})$ time, this method does too. □

Claim: The `set(idx)` method works in $O(\sqrt{n})$ time.

Proof.

The method consists of a single call to `getNode(idx)` and a value update. Since `getNode(idx)` works in $O(\sqrt{n})$ time, this method does too. □

Claim: The `insert(val, idx)` method works in amortized $O(\sqrt{n})$ time.

Proof.

The method calls `getNode(idx)` which is $O(\sqrt{n})$ and then creates a new node (which takes constant time with respect to the length of the list).

In every call it will make a call to `retreat(insertionIdx)`, which traverses the entire secondary list and thus takes $O(\sqrt{m})$ time.

Once every \sqrt{m} calls it will make a call to `buildMeta()` which traverses the entire main list so it takes $O(n)$ time. Distributing the cost of that call over the \sqrt{m} calls yields $O(\frac{n}{\sqrt{m}})$ time.

Therefore, overall, `insert(val, idx)` runs in $O(\sqrt{n}) + O(\sqrt{m}) + O(\frac{n}{\sqrt{m}})$ time, which is bounded by $O(\sqrt{n})$. □

Claim: The `remove(idx)` method works in amortized $O(\sqrt{n})$ time.

Proof.

This method works identically with the `insert(val, idx)` method, and therefore has the same runtime. □